Comparative Analysis of Cutting-Edge Transformer Models in Abstractive Text Summarization

Dhruva Sandu*
Department of Computer Engineering and Technology
Dr.Vishwanath Karad MIT WORLD PEACE UNIVERSITY
Pune, India
dhruvasandu@gmail.com

Pradnya Kulkarni
Department of Computer Engineering and Technology
Dr.Vishwanath Karad MIT WORLD PEACE UNIVERSITY
Pune, India
pradnya.kulkarni@mitwpu.edu.in

Ramaa Sandu
Department of Mathematics and Statistics
Dr.Vishwanath Karad MIT WORLD PEACE UNIVERSITY
Pune, India
ramaa.sandu@mitwpu.edu.in

***Abstract*—In today's information-rich landscape, the deluge of textual data poses a significant challenge, necessitating efficient methods for distilling key information. Text summarization emerges as a pivotal solution within the realm of Natural Language Processing (NLP), enabling the extraction of essential insights from copious volumes of text. This paper delves into the domain of abstractive text summarization, where the focus lies not only on selecting and rearranging existing text segments but also on generating novel phrases and sentences to convey the essence of the source material. The study evaluates the efficacy of advanced transformer-based models, including PEGASUS, BART, and T5, in the context of abstractive summarization. This research adopts the BERT score metric in addition to the traditional ROUGE metric to provide a comprehensive analysis of model performance, encompassing precision, recall, and F1 score. Leveraging two datasets comprising news articles sourced from the British Broadcasting Corporation (BBC) and CNN/Daily Mail, the models are evaluated to assess their ability to generate high-quality summaries. The results highlight PEGASUS as the front runner in terms of BERT Score metrics on the BBC News Dataset and the ROUGE Scores on the CNN/Daily Mail Dataset, with precision, recall, and F1 scores of 0.9246, 0.8584, and 0.8901, respectively. At the same time, T5 slightly edges PEGASUS and BART in terms of BERT Scores on the CNN/Daily Mail Dataset, with precision, recall, and F1 scores of 0.8733, 0.8631, and 0.8681, respectively. Overall, all models display commendable performance in terms of ROUGE and BERT Score on both datasets. Insights gleaned from the study not only shed light on the comparative performance of these models but also pave the way for future research endeavors. By pushing the boundaries of summarization technology, this research contributes to the advancement of NLP and fosters a deeper understanding of text summarization's practical utility across diverse domains.**

***Index Terms*—Text summarization, abstractive summarization, PEGASUS, BART, T5, deep learning, BERT score, ROUGE score NLP**

I. Introduction

Text summarization stands as a pivotal facet within the realm of natural language processing (NLP), offering a means to distill copious volumes of text into succinct and informative summaries. In today's information-rich landscape, the ability to effectively condense textual content holds paramount importance, enabling individuals and organizations to glean insights, make informed decisions, and navigate the deluge of information inundating digital platforms. Through the synthesis of key information, while preserving contextual nuances, text summarization serves as an invaluable tool across diverse domains, ranging from journalism and academia to business intelligence and legal analysis.

In order to delve deeper into the realm of text summarization and its implications, this paper is structured as follows: Section 1 introduces text summarization, emphasizing its importance and focusing on abstractive methods. It discusses the motivation for summarization, the limitations of extractive approaches, and the advantages of abstractive summarization. The section concludes with the research objective to evaluate transformer-based models, using the BERT score metric for a comprehensive performance analysis. Section 2 delves into the background and related work, offering an overview of existing approaches to text summarization and situating our study within the broader context of Natural Language Processing (NLP) research. Section 3 describes the methodology employed, detailing the datasets used, the architecture of abstractive text summarization, the transformer-based models selected for evaluation, and the specific metrics for performance assessment. Section 4 presents the experimental results, including a comparative analysis of PEGASUS, BART, and T5 models based on precision, recall, and F1 score. Section 5 discusses the implications of the findings, addressing the strengths and limitations of each model. Finally, Section 6 concludes the paper, summarizing the key insights and suggesting potential directions for future research in the field of abstractive text summarization.

A. Motivation for Text Summarization

The motivation for text summarization stems from the pervasive challenge of information overload, wherein the sheer abundance of textual data renders conventional methods of comprehension and analysis inadequate. In response, text summarization emerges as a strategic solution, offering a mechanism to distill relevant information, eliminate extraneous details, and present cohesive summaries that encapsulate the essence of the original content. Whether facilitating information retrieval, enabling content summarization, or supporting document clustering, text summarization facilitates enhanced productivity, decision-making, and knowledge discovery across myriad applications.

B. Extractive Text Summarization

Extractive text summarization involves selecting and concatenating existing sentences or phrases from the source text to create a summary [1]. However, this approach has notable drawbacks: it often results in summaries that lack cohesion and coherence, as the extracted sentences may not seamlessly fit together. Additionally, it tends to include redundant information by selecting multiple sentences conveying similar points, and it lacks the ability to rephrase or paraphrase, limiting the presentation of information in a concise or contextually appropriate manner. Extractive methods may fail to capture the underlying meaning or nuances of the original text, as they are constrained to selecting literal segments without understanding the context. Due to these limitations, there is a growing need for abstractive text summarization.

C. Abstractive Text Summarization

Of particular interest within the domain of text summarization is the paradigm of abstractive summarization. Unlike extractive summarization, which merely selects and concatenates existing text segments, abstractive summarization involves the synthesis of new phrases and sentences to convey the main ideas of the source text [1]. This approach offers greater flexibility and creativity in summarization, allowing for the generation of concise and coherent summaries that capture the underlying meaning and intent of the original content. However, abstractive summarization poses unique challenges, including the need for models to comprehend and generate natural language fluently while preserving the fidelity and coherence of the summary [1], [2], [4], [5].

D. Objective of the Research

This research aims to explore and evaluate the efficacy of advanced transformer-based models, including PEGASUS, BART, and T5, in the domain of abstractive text summarization. While prior research predominantly relies on metrics such as ROUGE score for evaluation, this study adopts the BERT score metric to provide a comprehensive analysis of summarization model performance. By leveraging state-of-the-art evaluation techniques, we seek to assess the strengths, limitations, and potential enhancements of these models in generating high-quality abstractive summaries. Through this endeavor, we aim to contribute to the advancement of text summarization technology and foster a deeper understanding of its practical utility in real-world applications.

II. Literature Review

In the realm of natural language processing (NLP), automatic text summarization has emerged as a crucial technique for condensing vast amounts of textual data into concise and informative summaries. Over the past decade, significant advancements have been made in both extractive and abstractive summarization approaches [1]–[3]. Extractive summarization involves selecting and rearranging existing sentences or phrases from the source text to create a summary, while abstractive summarization generates summaries by paraphrasing and rephrasing the original text to convey the same meaning using different words [1].

Recent years have witnessed a notable shift towards abstractive methods, driven by advancements in deep learning and transformer-based models [1-4]. Transformer models, known for their effectiveness in handling sequential data, have revolutionized the field of NLP [11]. These models, such as BART, T5, and PEGASUS, have demonstrated remarkable performance in various text summarization tasks such as news article summarization, document summarization, and conversational summarization [4], [12]–[14].

Among these models, BART, stands out for its denoising autoencoder architecture designed for pretraining sequence-to-sequence models [13]. This flexibility in applying various noising approaches to input text has contributed to its effectiveness in both text generation and comprehension tasks. However, BART's training process can be computationally intensive, requiring significant computational resources [13].

Similarly, PEGASUS introduces a novel self-supervised pretraining objective tailored for abstractive text summarization, achieving state-of-the-art performance across diverse summarization tasks [14]. However, one limitation of PEGASUS is its computational complexity, which can pose challenges for real-time applications. Additionally, PEGASUS may struggle with capturing fine-grained details and nuances in the text, leading to potentially oversimplified summaries [14].

Transfer learning has significantly advanced text summarization techniques, as highlighted by the Unified Text-to-Text Transformer framework [12]. This framework proposes a systematic approach to studying transfer learning across various NLP tasks, emphasizing the necessity for models to develop general-purpose knowledge. Leveraging transfer learning has led to recent advancements and state-of-the-art results in various NLP benchmarks, further underscoring the importance of this technique [12]. T5, a transfer learning model, demonstrates its versatility by performing various NLP tasks, including text summarization, with a single unified architecture [4]. This unified approach streamlines model deployment and maintenance processes. However, T5 may encounter challenges in generating diverse and fluent summaries, especially for longer texts [12].

Research has also explored the efficacy of different neural network architectures, including Long Short-Term Memory (LSTM), Gated Recurrent Unit (GRU), Recurrent Neural Network (RNN), and Transformer, for automatic text summarization [15]. These studies have highlighted the importance of transformer models, particularly in achieving high accuracy in summarization tasks [15].

Moreover, research has explored specialized applications of text summarization, such as personalized multi-document summarization and summarization tailored for specific domains like behavioral biology literature [8], [9]. These studies highlight the importance of considering user preferences and domain-specific requirements in the summarization process.

The studies reviewed collectively contribute to a deeper understanding of text summarization research trends and provide valuable guidance for future research directions.

III. Materials and Methods

A. Working of Abstractive Text Summarization

The Figure 1 presents the architecture of abstractive text summarization with a specific focus on transformer models.

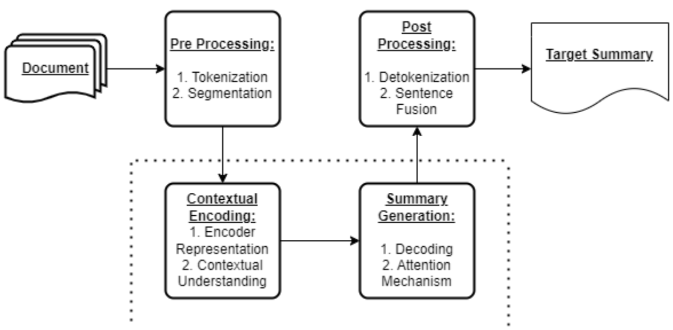

Fig. 1. Architecture of Abstractive Text Summarization

We will now explore each component of the architecture in more detail.

1) Pre-Processing: This step prepares the text for the summarization process. It involves tokenization, which breaks the text down into its basic units, usually words or punctuation marks. Segmentation can further divide the text into smaller units, such as sentences or paragraphs.

2) Contextual Encoding: This step captures the meaning of the text using transformer models. Encoder representation involves using a transformer encoder to process the text and generate a numerical representation that captures the meaning of the text. Contextual understanding incorporates additional information, such as word relationships or sentiment, to enrich the understanding of the text.

3) Summary Generation: This part of the model uses the encoded representation of the text to generate a summary. Decoding involves a transformer decoder, which uses the encoded representation to generate the summary one word at a time. The attention mechanism helps the decoder to focus on the most important parts of the encoded representation when generating the summary.

4) Post-Processing: This step refines the summary generated by the model. Detokenization converts the summary back into human-readable text by joining the tokens together. Sentence fusion might involve merging short and choppy sentences into a more grammatically sound summary.

B. Datasets

1) BBC News Summary Dataset: One of the dataset used for this research was sourced from the British Broadcasting Corporation (BBC) [https://www.kaggle.com/datasets/pariza/bbc-news-summary]. It comprises 2,225 news articles spanning 2004 to 2005. Alongside the articles are concise summaries, providing key information from the original text. These summaries are valuable for text summarization tasks.

The Figure 2 displays a sample news article and its corresponding reference summary from the dataset.

Microsoft seeking spyware trojan Microsoft is investigating a trojan program that attempts to switch off the firm's anti-spyware software. The spyware tool was only released by Microsoft in the last few weeks and has been downloaded by six million people. Stephen Toulouse, a security manager at Microsoft, said the malicious program was called Bankash-A Trojan and was being sent as an e-mail attachment. Microsoft said it did not believe the program was widespread and recommended users to use an anti-virus program. The program attempts to disable or delete Microsoft's anti-spyware tool and suppress warning messages given to users. It may also try to steal online banking passwords or other personal information by tracking users' keystrokes. Microsoft said in a statement it is investigating what it called a criminal attack on its software. Earlier this week, Microsoft said it would buy anti-virus software maker Sybari Software to improve its security in its Windows and e-mail software. Microsoft has said it plans to offer its own paid-for anti-virus software but it has not yet set a date for its release. The anti-spyware program being targeted is currently only in beta form and aims to help users find and remove spyware - programs which monitor internet use, causes advert pop-ups and slow a PC's performance.

Microsoft is investigating a trojan program that attempts to switch off the firm's anti-spyware software.Microsoft said it did not believe the program was widespread and recommended users to use an anti-virus program.Stephen Toulouse, a security manager at Microsoft, said the malicious program was called Bankash-A Trojan and was being sent as an e-mail attachment.Microsoft said in a statement it is investigating what it called a criminal attack on its software.

Fig. 2. Sample News Article and corresponding reference summary from the dataset

The news articles cover a diverse range of topics, including politics, business, sport, technology, and entertainment, reflecting the breadth and depth of news coverage provided by the BBC during the specified timeframe. Each news article is meticulously curated and categorized to ensure relevance and coherence within the dataset.

The dataset is organized into two main attributes: the ”News Articles” attribute contains the full-text articles, while the ”Summaries” attribute houses the corresponding summaries for each article.

2) CNN/Daily Mail Dataset: The second dataset utilized in this research is the CNN/Daily Mail Dataset [16]. Specifically, the study employed the test dataset for analysis. This dataset comprises 11,490 distinct news articles sourced from CNN/Daily Mail, covering a wide range of topics. Accompanying each article are reference summaries, which serve as valuable resources for text summarization tasks.

The Figure 3 showcases a representative example from the dataset, presenting an article alongside its corresponding highlight, providing insight into the content and structure of the dataset.

(CNN)Blinky and Pinky on the Champs Elysees? Inky and Clyde running down Broadway? Power pellets on the Embarcadero? Leave it to Google to make April Fools' Day into throwback fun by combining Google Maps with Pac-Man. The massive tech company is known for its impish April Fools' Day pranks, and Google Maps has been at the center of a few, including a Pokemon Challenge and a treasure map. This year the company was a day early to the party, rolling out the Pac-Man game Tuesday. It's easy to play: Simply pull up Google Maps on your desktop browser, click on the Pac-Man icon on the lower left, and your map suddenly becomes a Pac-Man course. Twitterers have been tickled by the possibilities, playing Pac-Man in Manhattan, on the University of Illinois quad, in central London and down crooked Lombard Street in San Francisco, among many locations.

Google Maps has a temporary Pac-Man function. Google has long been fond of April Fools' Day pranks and games. Many people are turning their cities into Pac-Man courses.

Fig. 3. Sample Article and corresponding highlight from the dataset

The dataset is structured into two primary attributes: ”article” and ”highlight.” The ”article” attribute contains the full-text news articles, while the ”highlight” attribute includes their corresponding summaries. This organization ensures ease of access and facilitates efficient retrieval of pertinent information for summarization purposes.

The news articles within this dataset span various domains, reflecting the extensive coverage provided by CNN/Daily Mail. The dataset encompasses a diverse array of topics, enabling comprehensive exploration and analysis.

C. Models

The Figure 4 illustrates the functioning of three transformer based models—PEGASUS-Large, BART-Large, and T5-Large—in the context of abstractive text summarization.

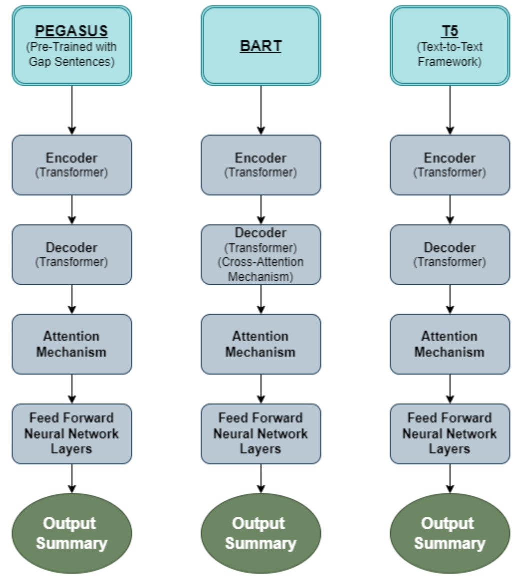

Fig. 4. Model Architectures for Abstractive Text Summarization

1) PEGASUS-Large: The PEGASUS-Large model, developed by Google Research, is the first transformer-based pretrained model used for this study. PEGASUS is tailored specifically for abstractive text summarization tasks. It is built upon a transformer architecture, which excels in capturing intricate linguistic patterns and generating coherent text sequences [14]. The ”Large” variant of PEGASUS represents an enhanced version optimized for processing large volumes of text data efficiently. Leveraging a diverse corpus of text from various sources such as news articles, Wikipedia, books, and websites, PEGASUS-Large has been trained to understand and summarize complex textual content effectively. Through its innovative self-supervised learning approach known as Pretraining with Extracted Gap-sentences for Abstractive Summarization Sequences (PEGASUS), this model reconstructs gapped-out sections of text to generate high-quality abstractive summaries that encapsulate the essence of the input text [14].

2) BART-Large: BART (Bidirectional and Auto Regressive Transformers), developed by Facebook AI, represents another cutting-edge transformer-based model designed for abstractive text summarization tasks. The base BART model combines bidirectional and auto-regressive capabilities, allowing it to understand and generate text sequences in both directions [13]. The ”Large” variant, BART-Large, denotes an extensive model trained on diverse text corpora, including news articles, books, and websites. By leveraging a training objective focused on denoising auto encoding, BART-Large learns to reconstruct the original, uncorrupted text from corrupted input text, enabling it to capture coherent representations of textual content effectively [13]. This is another transformer-based pre-trained model utilized in this study.

3) T5-Large: The T5-Large model, developed by Google Research, is the last transformer-based pre-trained model used for this study. T5 (Text-To-Text Transfer Transformer) is capable of addressing a broad spectrum of text generation tasks, including abstractive text summarization. The foundational T5 model follows a text-to-text framework, wherein both input and output are represented as textual sequences. The "Large" variant, T5-Large, signifies an augmented model trained on a diverse array of text data encompassing news articles, books, and websites [12]. By exposing T5-Large to various tasks such as summarization, translation, and question answering within a unified text-to-text format during training, this model learns to map input text to output text across disparate tasks, thereby demonstrating remarkable adaptability and performance in abstractive text summarization scenarios.

D. Evaluation Method

To provide a more comprehensive understanding of the evaluation metrics used in this research project, it's valuable to delve deeper into the comparison between BERT score and ROUGE.

- ROUGE (Recall Oriented Understudy for Gisting Evaluation) is a traditional metric widely used for evaluating the quality of automatic text summarization systems. It primarily relies on n-gram overlap between the generated summary and reference summaries to compute precision, recall, and F1 score. While ROUGE has been a cornerstone in the field, it comes with certain limitations. One key drawback is its reliance on exact word matches, which may not effectively capture the semantic similarity between the generated and reference summaries. Additionally, ROUGE tends to overlook syntactic variations and paraphrasing, potentially penalizing well-written summaries that convey the same information using different wording or sentence structures.
- In contrast, BERT(Bidirectional Encoder Representations from Transformers) score represents a more contemporary and advanced evaluation metric, leveraging contextualized embeddings derived from pre-trained BERT models. Unlike ROUGE, which focuses solely on surface-level textual overlap, BERT score captures the semantic similarity between the generated and reference summaries by considering the contextual embeddings of words and phrases. This enables BERT score to provide a more nuanced evaluation, taking into account not only exact word matches but also semantic equivalence and syntactic variations. By incorporating contextual embeddings, BERT score offers a more robust assessment of summary quality, which aligns more closely with human judgment. BERT score provides values for precision, recall, and F1-score, offering a comprehensive analysis of summary quality.

1) Precision: Precision measures the proportion of the summary generated by the model that is relevant to the reference summary. Precision indicates how well the generated summary captures important information present in the reference summary.

$$\text{Precision} = \frac{\text{Number of relevant words in the generated summary}}{\text{Total number of words in the generated summary}} \tag{1}$$

2) Recall: Recall measures the proportion of relevant information from the reference summary that is successfully captured by the generated summary. In text summarization, recall reflects the coverage of essential details from the reference summary in the generated summary.

$$\text{Recall} = \frac{\text{Total number of relevant words in the reference summary}}{\text{Number of words in the generated summary}} \tag{2}$$

3) F1-Score: F1 score is the harmonic mean of precision and recall, providing a balanced assessment of both metrics. It serves as a single scalar value representing the overall quality of the generated summary compared to the reference summary.

$$\text{F1-Score} = 2\left(\frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}\right) \tag{3}$$

IV. Results and Discussions

A. Actual Summary Examples

The subsequent Figure 5 demonstrates summaries produced by each of the three models in response to the news article depicted in Figure 2.

Microsoft seeking spyware trojan Microsoft is investigating a trojan program that attempts to switch off the firm's anti-spyware software. Earlier this week, Microsoft said it would buy anti-virus software maker Sybari Software to improve its security in its Windows and e-mail software.

Microsoft is investigating a trojan program that attempts to switch off the firm's anti-spyware software. The spyware tool was only released by Microsoft in the last few weeks and has been downloaded by six million people. Microsoft said it did not believe the program was widespread.

the spyware tool was only released by Microsoft in the last few weeks. it has been downloaded by six million people. Microsoft said it did not believe the program was widespread.

Fig. 5. Generated Summary examples of PEGASUS-Large, BART-CNN-Large and T5-Large

In comparing the summaries generated by PEGASUS-Large, BART-Large, and T5-Large for the news article provided in Figure 2, it's evident that each model captures the essence of the original text with varying degrees of detail and efficiency. BART-Large stands out for its comprehensive summary, efficiently condensing key information while maintaining clarity. Notably, it was also the fastest in producing the summary. Despite minor punctuation issues, T5-Large provides a succinct overview of the article. PEGASUS-Large offers a concise summary but lacks some detail.

B. Evaluation on BBC News Dataset

First, we examine the evaluation of each of the three models on the BBC News dataset based on BERT scores. For this evaluation, 100 random samples were selected from the dataset. The results of the evaluation, showcasing the average values of precision, recall, and F1-score across all three models, are presented in Table I.

TABLE I
EVALUATION OF MODELS ON BBC NEWS DATASET

Models	**BERT Score**		
	Precision	**Recall**	**F1-Score**
PEGASUS-Large	0.9246	0.8584	0.8901
BART-Large	0.8708	0.8543	0.8619
T5-Large	0.8708	0.8543	0.8619

The evaluation of the models on the BBC News Dataset based on BERT scores unveils intriguing observations regarding their performance in abstractive text summarization. PEGASUS-Large emerges as the frontrunner, showcasing exceptional precision, recall, and F1-Score metrics. This superior performance underscores the model's adeptness in generating summaries that closely mirror the essence of the source text. The robustness of PEGASUS-Large highlights its capability to capture nuanced information and produce coherent summaries, making it a compelling choice for tasks requiring high-quality abstractive summarization.

In contrast, BART-Large and T5-Large demonstrate comparable performance, albeit with slightly lower scores than PEGASUS-Large. While both models exhibit commendable precision, recall, and F1-Score metrics, they fall short of PEGASUS-Large in terms of overall effectiveness. This marginal difference in performance suggests that BART-Large and T5-Large are still viable options for abstractive summarization tasks, particularly when considering factors such as computational resources and model complexity.

C. Evaluation on CNN/Daily Mail Dataset

Now, we delve into the evaluation of the models on the CNN/Daily Mail Dataset. To conduct this evaluation, 100 random samples were drawn from the dataset. Table II presents the results, encompassing the average values of precision, recall, and F1-score, as determined by the BERT score metric.

TABLE II
EVALUATION OF MODELS ON CNN/DAILY MAIL DATASET

Models	**BERT Score**		
	Precision	**Recall**	**F1-Score**
PEGASUS-Large	0.8552	0.8587	0.8568
BART-Large	0.8591	0.8738	0.8663
T5-Large	0.8733	0.8631	0.8681

PEGASUS-Large demonstrates a high precision score of 0.8552, indicating its ability to generate summaries that are highly relevant to the reference summaries. However, its recall and F1-Score are slightly lower, suggesting that it may miss some important details from the source text. BART-Large performs well across all metrics, with a precision of 0.8591, recall of 0.8738, and F1-Score of 0.8663. This indicates that BART-Large is effective in both generating relevant summaries and capturing key information from the source text. T5-Large also performs well, with a precision of 0.8733, recall of 0.8631, and F1-Score of 0.8681. This indicates that T5-Large is able to generate summaries that are both relevant and comprehensive.

In a related study by [14], the same models were evaluated on the CNN/Daily Mail Dataset using the ROUGE metric. The results of this evaluation are provided in Table III.

TABLE III
ROUGE SCORE EVALUATION ON CNN/DAILY MAIL DATASET

Models	**ROUGE Score**		
	R1	**R2**	**RL**
PEGASUS-Large	44.17	21.47	41.11
BART-Large	44.16	21.28	40.9
T5-Large	43.52	21.55	40.69

In Table III, all three models are compared using the ROUGE Score metric. The three main ROUGE metrics are R1, R2, and RL. R1 measures the overlap of unigrams (single words), R2 measures the overlap of bigrams (pairs of adjacent words), and RL measures the longest common subsequence.

Through the evaluation, we can gather that PEGASUS-Large stands out with the highest scores across all three ROUGE metrics: R1 (44.17), R2 (21.47), and RL (41.11). This indicates that PEGASUS-Large generates summaries that closely match the reference summaries in terms of both unigram and bigram overlap, as well as capturing the longest common subsequence. While BART-Large and T5-Large also perform well, with slightly lower scores, PEGASUS-Large demonstrates superior performance in content overlap.

Combining the evaluations from both BERT Score and ROUGE Score metrics on the CNN/Daily Mail dataset provides a comprehensive analysis of the performance of PEGASUS-Large, BART-Large, and T5-Large models. According to the BERT Score evaluation, T5-Large slightly outperforms the other models in terms of F1-Score, indicating its ability to generate relevant and comprehensive summaries. However, PEGASUS-Large demonstrates high precision, suggesting that its generated summaries are highly relevant to the reference summaries, although it lags slightly in recall and F1-Score. On the other hand, the ROUGE Score evaluation highlights PEGASUS-Large as the superior model with the highest scores in R1, R2, and RL, reflecting its strong performance in capturing both unigram and bigram overlaps and the longest common subsequence with the reference summaries. BART-Large and T5-Large also perform well but show slightly lower scores in ROUGE metrics. Overall, PEGASUS-Large appears to excel in content overlap and relevance according to the ROUGE scores, while T5-Large shows a balanced performance in relevance and comprehensiveness as per the BERT Score. This dual-metric analysis underscores the importance of employing multiple evaluation criteria to fully understand the strengths and limitations of these transformer-based models in abstractive summarization tasks.

V. Conclusion and Future Work

This paper investigates the efficacy of transformer-based models—PEGASUS-Large, BART-Large, and T5-Large—for abstractive text summarization tasks on the BBC News and CNN/Daily Mail datasets. Comprehensive evaluations using both BERT Score and ROUGE metrics reveal the models' ability to generate coherent and informative summaries. PEGASUS-Large outperforms other models on the BBC News dataset in terms of BERT Score, demonstrating superior precision, recall, and F1-Score metrics. BART-Large and T5-Large also perform commendably on the same dataset, though slightly behind PEGASUS-Large in BERT Score. For the CNN/Daily Mail dataset, PEGASUS-Large achieves the highest ROUGE scores, indicating its effectiveness in capturing content overlap with reference summaries. However, T5-Large shows slightly better performance in BERT Score metrics compared to PEGASUS-Large and BART-Large. The comparison between BERT Score and ROUGE metrics provides valuable insights into the strengths and limitations of each model, underscoring the importance of employing multiple evaluation criteria for a comprehensive assessment.

In addition to the promising results obtained, future works could explore innovative approaches such as ensemble methods to combine the strengths of different models, domain-specific fine-tuning to enhance performance on specific types of content, and incorporating multi-modal inputs to leverage additional contextual information for summarization tasks. These approaches have the potential to further advance the capabilities of transformer-based models and address challenges in abstractive text summarization across diverse domains and modalities.

Overall, the results indicate promising prospects for transformer-based models in abstractive text summarization, warranting further investigation and refinement of these models for real-world applications.

References

[1] A. P. Widyassari et al., "Review of automatic text summarization techniques amp; methods," Journal of King Saud University - Computer and Information Sciences, vol. 34, no. 4. Elsevier BV, pp. 1029–1046, Apr. 2022.

[2] D. Yadav, J. Desai, and A. K. Yadav, "Automatic Text Summarization Methods: A Comprehensive Review." arXiv, 2022.

[3] S. Alhojely and J. Kalita, "Recent Progress on Text Summarization," 2020 International Conference on Computational Science and Computational Intelligence (CSCI). IEEE, Dec. 2020.

[4] Ö. B. Mercan, S. N. Cavsak, A. Deliahmetoglu, and S. Tanberk, "Abstractive Text Summarization for Resumes With Cutting Edge NLP Transformers and LSTM." arXiv, 2023.

[5] N. K. Sirohi, Dr. M. Bansal, and Dr. S. N. R. Rajan, "Text Summarization Approaches Using Machine Learning amp; LSTM," Revista Gestão Inovação e Tecnologias, vol. 11, no. 4. Centivens Institute of Innovative Research, pp. 5010–5026, Sep. 01, 2021.

[6] S.-N. Vo, T.-T. Vo, and B. Le, "Interpretable extractive text summarization with meta-learning and BI-LSTM: A study of meta learning and explainability techniques," Expert Systems with Applications, vol. 245. Elsevier BV, p. 123045, Jul. 2024.

[7] M. Mohd et al., "Semantic-Summarizer: Semantics-based text summarizer for English language text," Software Impacts, vol. 18. Elsevier BV, p. 100582, Nov. 2023.

[8] S. Chaurasia, D. Dasgupta, and R. Regunathan, "T5LSTM-RNN based Text Summarization Model for Behavioral Biology Literature," Procedia Computer Science, vol. 218. Elsevier BV, pp. 585–593, 2023.

[9] K. Veningston, P. V. Venkateswara Rao, and M. Ronalda, "Personalized Multi-document Text Summarization using Deep Learning Techniques," Procedia Computer Science, vol. 218. Elsevier BV, pp. 1220–1228, 2023.

[10] G. Van Houdt, C. Mosquera, and G. Nápoles, "A review on the long short-term memory model," Artificial Intelligence Review, vol. 53, no. 8. Springer Science and Business Media LLC, pp. 5929–5955, May 13, 2020.

[11] T. T, W. Haryono, A. U. Zailani, Y. M. Djaksana, N. Rosmawarni, and N. D. Arianti, "Transformers in Machine Learning: Literature Review," Jurnal Penelitian Pendidikan IPA, vol. 9, no. 9. Universitas Mataram, pp. 604–610, Sep. 25, 2023.

[12] C. Raffel et al., "Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer." arXiv, 2019.

[13] M. Lewis et al., "BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension." arXiv, 2019.

[14] J. Zhang, Y. Zhao, M. Saleh, and P. J. Liu, "PEGASUS: Pre-training with Extracted Gap-sentences for Abstractive Summarization." arXiv, 2019.

[15] S. BAYAT and G. ISIK, "ASSESSING THE EFFICACY OF LSTM, TRANSFORMER, AND RNN ARCHITECTURES IN TEXT SUMMARIZATION," International Conference on Applied Engineering and Natural Sciences, vol. 1, no. 1. All Sciences Proceedings, pp. 813–820, Jul. 21, 2023.

[16] A. See, P. J. Liu, and C. D. Manning, "Get To The Point: Summarization with Pointer-Generator Networks," Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Association for Computational Linguistics, 2017.